

Assignment9 Report

姓名: Yitong WANG(王奕童) 11910104@mail.sustech.edu.cn

学号: 11910104

实验课时段: 周五5-6节

实验课教师: Yun SHEN(沈昀) sheny@mail.sustech.edu.cn

实验课SA:

- Yining TANG(汤怡宁) 11811237@mail.sustech.edu.cn
- Yushan WANG(王宇杉) 11813002@mail.sustech.edu.cn

1. Disk scheduling

(1) READ/WRITE data time = Seek Time + Rotational Latency + Transfer Time

(2)

a. Track Access Sequence:

- FIFO: 70 => 30 => 90 => 120 => 60 => 20
- SSTF: 90 => 70 => 60 => 30 => 20 => 120
- SCAN: 120 => 90 => 70 => 60 => 30 => 20
- CSCAN: 120 => 20 => 30 => 60 => 70 => 90

b. Time Calculation

For the seek time, first calculate their tracks:

- FIFO: $(100-70) + (70-30) + (90-30) + (120-90) + (120-60) + (60-20) = 260$
- SSTF: $(100-90) + (90-70) + (70-60) + (60-30) + (30-20) + (120-20) = 180$
- SCAN: $(120-100) + (199-120) + (199-90) + (90-70) + (70-60) + (60-30) + (30-20) = 279$
- CSCAN: $(120-100) + (199-120) + (199-0) + (20-0) + (30-20) + (60-30) + (70-60) + (90-70) = 388$

Thus time is:

- FIFO: 260ms

- SSTF: 180ms
- SCAN: 279ms
- CSCAN: 388ms

For FIFO\SSTF\SCAN\CSCAN algorithm, their rotational latency is the same.

12000r/min => 200r/s => 0.2r/ms => 5ms/r

Since it is randomly distributed access, we treat it as half round: 2.5ms

Since 6 accesses, total time = 2.5ms*6 = 15ms

The question does not tell us the transfer time, thus we omitted.

Total Time:

- FIFO: 260+15=275ms
- SSTF: 180+15=195ms
- SCAN: 279+15=294ms
- CSCAN: 388+15=403ms

2. Simple File System

Consider adding one statement in `tools/mksfs.c` :

```

428 static void
429 add_entry(struct sfs_fs *sfs, struct cache_inode *current, struct cache_inode *file, const char *name) {
430     static struct sfs_entry __entry, *entry = &__entry;
431     assert(current->inode.type == SFS_TYPE_DIR && strlen(name) <= SFS_MAX_FNAME_LEN);
432     entry->ino = file->ino, strcpy(entry->name, name);
433     uint32_t entry_ino = sfs_alloc_ino(sfs);
434     write_block(sfs, entry, sizeof(struct sfs_entry), entry_ino);
435     append_block(sfs, current, sizeof(struct sfs_entry), entry_ino, name);
436     printf("entry is %d, name is %s, file->ino is %d\n", entry_ino, name, file->ino);
437     file->inode.nlinks ++;
438 }
439 }
440

```

And the partial output of `make qemu` is:

```
48016 records out
480000 bytes (480 kB, 469 KiB) copied, 0.00130072 s, 369 MB/s
entry is 3, name is ., file->ino is 1
entry is 4, name is .., file->ino is 1
entry is 7, name is test, file->ino is 5
entry is 12, name is hello, file->ino is 8
entry is 17, name is sh, file->ino is 13
entry is 19, name is ., file->ino is 18
entry is 20, name is .., file->ino is 1
entry is 23, name is file1, file->ino is 21
entry is 24, name is dir1, file->ino is 18
create bin/sfs.img (disk0) successfully.
+ ld bin/kernel
riscv64-unknown-elf-objcopy bin/kernel --strip-all -O binary bin/ucore.img
```

Thus we draw the diagram according to the running result:

| | | | | |
|--------|----------------|--|---|----|
| 0 | superblock | | 0 | 3 |
| 1 | rootdir | | 1 | 4 |
| inode | freemap | | 2 | 7 |
| 3 | name='.' | | 3 | 12 |
| entry | ino=1 | | 4 | 17 |
| 4 | name='..' | | 5 | 24 |
| entry | ino=1 | | | |
| 5 | test文件的inode | | 0 | 6 |
| inode | test data | | | |
| 6 | name='test' | | | |
| entry | ino=5 | | | |
| 8 | hello文件的inode | | 0 | 9 |
| inode | hello data 1 | | 1 | 10 |
| 9 | hello data 2 | | 2 | 11 |
| 10 | hello data 3 | | | |
| 11 | name='hello' | | | |
| entry | ino=8 | | | |
| 12 | sh文件的inode | | 0 | 14 |
| inode | sh data 1 | | 1 | 15 |
| 13 | sh data 2 | | 2 | 16 |
| 14 | sh data 3 | | | |
| 15 | name='sh' | | 0 | 19 |
| entry | ino=13 | | 1 | 20 |
| 16 | dir1的inode | | 2 | 23 |
| inode | name='.'(一个点) | | | |
| 17 | ino=18 | | | |
| 18 | name='..'(两个点) | | | |
| entry | ino=1 | | | |
| 20 | file1的inode | | 0 | 22 |
| inode | file1 data | | | |
| 21 | name='file1' | | | |
| entry | ino=21 | | | |
| 22 | name='dir1' | | | |
| entry | ino=18 | | | |
| 25-116 | 空闲状态 | | | |